

# Deploying Main Application

Estimated time: 90 minutes

Welcome to the lab **Deploying Main Application**. You should have all the code working from the previous modules in this capstone. The lab focuses on deploying the **Django** application on IBM Kubernetes Services.

## Learning Objectives:

After completing this lab you will be able to:

1. Create a Dockerfile to containerize Django applications
2. Use the kubectl CLI
3. Work with Kubernetes deployment yaml file
4. Create a Kubernetes Deployment on IKS

**Note:** Kindly complete the lab in a single session without any break because the lab may go on offline mode and may cause errors. If you face any issues/errors during the lab process, please log out from the lab environment. Then clear your system cache and cookies and try to complete the lab.

## Verify the environment and command line tools

1. If a terminal is not already open, open a terminal window by using the menu in the editor: Terminal > New Terminal.

**Note:** Please skip this step if the terminal already appears.

2. Verify that kubectl CLI is installed.

```
kubectl version
```

3. Change to your project folder.

```
cd /home/project
```

## Exercise 1: Django - Clone the Github Repository and Set Up Environment

You created a new repository for the Django application from the provided template in a previous lab. If not, go back to the [Create Band Website with Django](#) lab and ensure you complete it before coming back to this lab.

1. Open a terminal with Terminal > New Terminal if one is not open already.
2. Next, use the export GITHUB\_ACCOUNT command to export an environment variable that contains the name of your GitHub account.

**Note:** Substitute your real GitHub account for the {your\_github\_account} placeholder below:

```
export GITHUB_ACCOUNT={your_github_account}
```

3. Then use the following commands to clone your repository.

```
git clone https://github.com/$GITHUB_ACCOUNT/Back-end-Development-Capstone.git
```

4. Change to the Back-end-Development-Capstone directory, and execute the `./bin/setup.sh` command.

```
cd Back-end-Development-Capstone  
bash ./bin/setup.sh
```

5. You should see the follow at the end of the setup execution:

6. Finally, use the `exit` command to close the current terminal. The environment will not be fully active until you open a new terminal in the next step.

## Exercise 2: Django - Connect Services and Run Locally

You need to have the `pictures` and the `songs` microservices running from the previous lab in order to complete this part. If you don't have them running please ensure you walk through the previous labs and take note of the service URLs to fill in here. Also ensure you are able to perform all the curl operations in the previous lab before continuing here.

Change to the server directory.

```
cd Back-end-Development-Capstone
```

Open `views.py` in the editor.

[Open `views.py` in IDE](#)

1. The `def songs(request):` method currently returns an array of songs. Change the code as follows and replace the `SONG_URL` with the songs microservice deployed in the previous lab.

```
def songs(request):  
    songs = req.get("SONGS_URL/song").json()  
    return render(request, "songs.html", {"songs": songs["songs"]})
```

2. The `def photos(request):` method currently returns an array of songs. Change the code as follows and replace the `PHOTO_URL` with the pictures microservice deployed in the previous lab.

```
def photos(request):  
    photos = req.get("PHOTO_URL/picture").json()  
    return render(request, "photos.html", {"photos": photos})
```

3. Perform migrations to create necessary tables.

```
python3 manage.py makemigrations
```

4. Run migration to activate models for the app.

```
python3 manage.py migrate
```

Take a screenshot of the terminal after executing the `migrate` command and save as `django-songs-photos-migrate.jpg` (or `.png`).

5. Start the local development server.

```
python3 manage.py runserver
```

If the application does not run locally, you should return to module 3 and ensure you have finished the final lab. The application should connect to the **songs** and **photos** services that are running on IBM Code Engine and RedHat OpenShift.

**Note:** Once you launch the application locally and navigate to the /song route, if the song lyrics are not displayed, please proceed with the Docker deployment and check if the song lyrics render correctly.

## Exercise 3: Django - Finish the Dockerfile

### Your Tasks

Unlike the other services, there is no option to deploy the main Django application using the `source-to-image` method to IBM Kubernetes Service. You will have to write a Dockerfile to build the image. Your task in this exercise is to fill out the incomplete Dockerfile.

Open the `Back-end-Development-Capstone/Dockerfile` file in the editor.

[Open Dockerfile in IDE](#)

1. Insert code to use a python base image.

Remember the way to specify a base image for your image is with the `FROM` statement in the Dockerfile. Use the `FROM` statement to specify `FROM python:3.9.16-slim` as the base image.

▼ Click here for a hint.

```
FROM python:3.8.2
```

2. Insert code to change the working directory.

Recall that `WORKDIR` is used in Dockerfile to set the working directory for any subsequent files.

▼ Click here for a hint.

```
WORKDIR $APP
```

### 3. Insert code to copy the requirements.txt file to \$APP

Recall that the `requirements.txt` text file lists all the Python dependencies required by a Django project to run. You need to copy it to the container.

▼ Click here for a hint.

```
COPY requirements.txt $APP
```

### 4. Next, you need to tell the container to install all the dependencies from the `requirements.txt` file using the `pip` tool using the `RUN` command syntax.

▼ Click here for a hint.

```
RUN pip3 install -r requirements.txt
```

### 5. Once the requirements have been installed, you can copy the rest of the source code to the image using the `COPY` command again.

▼ Click here for a hint.

```
COPY . $APP
```

### 6. The next step is to expose port 8000 using the `EXPOSE` command. This command informs Docker that the container will listen on the specified network ports at runtime. It does not actually publish the port to the host machine but rather provides metadata to help users understand how to interact with the container.

▼ Click here for a hint.

```
EXPOSE 8000
```

### 7. Great! You are almost done. The last line in the Dockerfile should tell Docker how to run the container. We want to run the `python manage.py runserver 0.0.0.0:8000` command when the container starts. Note the `0.0.0.0:8000` at the end. Without the `0.0.0.0:8000` argument, the server would only listen on the loopback interface (127.0.0.1) and would not be accessible from outside the container.

▼ Click here for a hint.

```
CMD ["python", "manage.py", "runserver", "0.0.0.0:8000"]
```

## Final Solution

Ensure your code looks like the this:

▼ Click here for a hint.

```
FROM python:3.9.16-slim
ENV PYTHONBUFFERED 1
ENV PYTHONWRITEBYTECODE 1
RUN apt-get update \
    && apt-get install -y netcat
ENV APP=/app
# Change the workdir.
WORKDIR $APP
# Install the requirements
COPY requirements.txt $APP
RUN pip3 install -r requirements.txt
# Copy the rest of the files
COPY . $APP
EXPOSE 8000
RUN chmod +x /app/entrypoint.sh
ENTRYPOINT ["/bin/bash","/app/entrypoint.sh"]
CMD ["python", "manage.py", "runserver", "0.0.0.0:8000"]
```

## Exercise 4: Django - Build and Upload Image to Registry

The next step is build the Docker image and then push it to the lab IBM Container Registry.

### Your Tasks

1. Build the Docker image using the `docker build` command.

▼ Click here for a hint.

```
docker build . -t djangoserver
```

Take a screenshot of the output of `docker build` command in the terminal and save as `deploy-getdjango-1.jpg` (or `.png`).

2. Export the `MY_NAMESPACE` variable. This is your namespace in the lab and the container registry. We will push the image to this namespace in the next step.

```
export MY_NAMESPACE=sn-labs-$USERNAME
```

You can echo the `MY_NAMESPACE` variable in the terminal to see the value. In my case, the variable is `sn-labs-captainfedor1`. It would be something different for you:

```
echo $MY_NAMESPACE
```

3. Tag the image with the correct path for your IBM Container Registry. Your container registry in the lab environment has the format `us.icr.io/$MY_NAMESPACE`.

▼ Click here for a hint.

```
docker tag djangoserver us.icr.io/$MY_NAMESPACE/djangoserver:1
```

Note that the `MY_NAMESPACE` variable is automatically populated from the environment if you correctly set it with the `export` command in the previous step.

4. Push the docker image to the ICR regisry:

```
docker push us.icr.io/$MY_NAMESPACE/djangoserver:1
```

This command may take a couple of minutes to complete. You should see progress in the terminal as follows:

```
The push refers to repository [us.icr.io/sn-labs-captainfedo1/djangoserver]
f82006359bac: Pushed
ca49210f87e2: Pushed
7db8dc861f85: Pushing [=====] 26.84MB/69.14MB
6308aca94bbb: Pushed
855c8cc5561a: Pushed
730529d9e57d: Pushed
550beda58888: Pushed
43b1c6b4f686: Pushed
7ca0a8cd8bdd: Pushed
36e5ac463d7d: Pushed
4695cdfb426a: Pushing [=====] 31.98MB/80.51MB
```

5. You can check that the image got pushed to your registry namespace with this command:

```
ibmcloud cr images --restrict $MY_NAMESPACE
```

You should see an output similar to the following. You might have images here from other labs as well.

Listing images...						
Repository	Tag	Digest	Namespace	Created	Size	Security status
us.icr.io/sn-labs-captainfedo1/djangoserver	1	41932716f7d5	sn-labs-captainfedo1	1 week ago	153 MB	-

## Evidence

1. Take a screenshot of the `docker build` command running in the lab terminal. Save the screenshot as `deploy-getdjango-1.jpg` (or `.png`).

## Exercise 5: Django - Complete the Deployment File

The lab template contains an incomplete `deployment.yaml` file. You will complete the file in this step.

Open the `Back-end-Development-Capstone/deployment.yaml` file in the editor.

[Open deployment.yaml in IDE](#)

1. Insert kind of Deployment in the code below:

```
kind: # {insert kind here}
```

2. Insert image name in the code below:

```
image: # {insert image here}
```

3. Insert port for the Django application in the code below:

```
containerPort: # {insert port here}
```

## Final Solution

Ensure your code looks like the following:

▼ Click here for a hint.

Replace `$MY_NAMESPACE` with your namespace.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    run: djangoserver
  name: djangoserver
spec:
  replicas: 1
  selector:
    matchLabels:
      run: djangoserver
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
  template:
    metadata:
      labels:
        run: djangoserver
    spec:
      containers:
        - image: us.icr.io/$MY_NAMESPACE/djangoserver:1
          imagePullPolicy: Always
          name: djangoserver
      ports:
```

```
- containerPort: 8000
  protocol: TCP
restartPolicy: Always
```

## Exercise 6: Django - Deploy Django application to IKS

1. In order to deploy the application to IBM Kubernetes Service, you can simply apply the deployment file:

```
kubectl apply -f ./deployment.yml
```

2. You can see the pods running by using the command:

```
kubectl get pods -w
```

The `-w` flag will watch the output. You should see the `djangoserver` pod in running status after going through other states:

```
$ kubectl get pods -w
NAME           READY   STATUS    RESTARTS   AGE
djangoserver-7d77df49f8-2g7kc  0/1     Pending   0          0s
djangoserver-7d77df49f8-2g7kc  0/1     Pending   0          1s
djangoserver-7d77df49f8-2g7kc  0/1     ContainerCreating  0          1s
djangoserver-7d77df49f8-2g7kc  0/1     ContainerCreating  0          2s
djangoserver-7d77df49f8-2g7kc  1/1     Running   0          9s
```

The `watch` flag will keep the `kubectl get` command open. You can use `ctrl+c` keys to close the command so the terminal is available for other commands.

3. Next, in order to access the deployment from outside the cluster, you need to use the `kubectl port-forward` command to forward the ports from the lab environment into the running pod. Replace the pod name with your pod name in the command below:

```
kubectl port-forward po/djangoserver-688dc44d47-bs5x7 8000:8000
```

You should see an output similar to this:

```
$ kubectl port-forward po/djangoserver-7d77df49f8-2g7kc 8000:8000
Forwarding from 127.0.0.1:8000 -> 8000
Forwarding from [::1]:8000 -> 8000
```

4. You can proceed to launch the application in a new tab by clicking `Launch Application` icon on the left bar. Once the tab opens, you can enter port as `8000` and click the `Your Application` button.

## Evidence

1. Take a screenshot of the `kubectl apply -f` command running in the lab terminal. Save the screenshot as `deploy-getdjango-3.jpg` (or `.png`).

## Evidence Checklist

Now that you have the application running in the lab environment, you need to take a few more screenshots to submit for peer review.

1. Take a screenshot named `deploy-getdjango-4.jpg` (or a `.png`) showing the main Django application running in the browser.
2. Take a screenshot named `deploy-getdjango-5.jpg` (or a `.png`) showing you are able to see the songs page.
3. Take a screenshot named `deploy-getdjango-6.jpg` (or a `.png`) showing you are able to see the pictures page.
4. Take a screenshot named `deploy-getdjango-7.jpg` (or a `.png`) showing you are able to log in as the admin user.
5. Take a screenshot named `deploy-getdjango-8.jpg` (or a `.png`) showing you are able to see concerts on the concert page. You will have to create one or more concerts as the admin user first.

Congratulations! You have completed the capstone project. You can proceed to submit your evidence for peer review and also review one of your peer's submission.

## Author(s)

CF

© IBM Corporation. All rights reserved.