

Configuring and Visualizing Metrics with Prometheus and Grafana



Estimated Time: 30 minutes

Welcome to the **Configuring and Visualizing Metrics with Prometheus and Grafana** lab. In this lab, you will become familiar with using Prometheus to monitor sample servers simulated with node exporters. Node exporters simulate server applications that serve metrics through the `/metrics` endpoint. You will use Prometheus to monitor the target `node_exporter` application by scraping its metrics endpoints, and you will also understand the Grafana UI.

Learning Objectives

After completing this exercise, you should be able to perform the following tasks:

- Create queries to get the metrics about the target
- Find the status of the targets
- Find some information about the targets and visualize it with graphs
- Create and customize graphs in Prometheus to visualize the metrics and status of monitored targets.
- Configure and launch Grafana to source data from Prometheus.

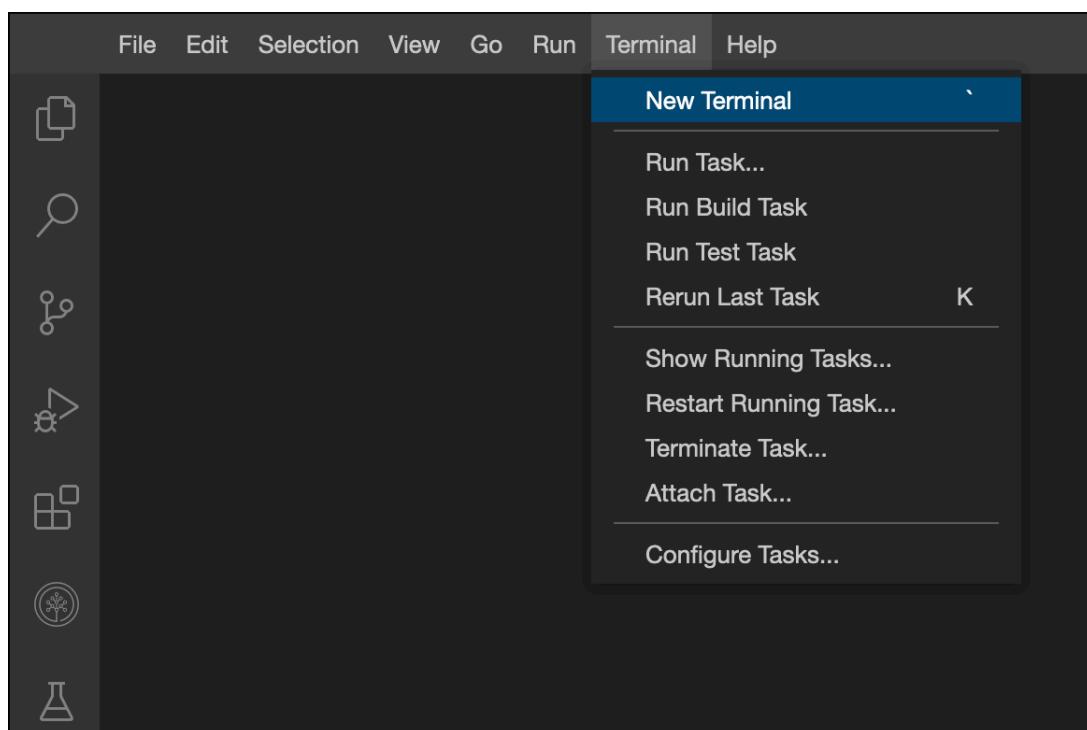
Prerequisites

Pre-requisites

This lab uses Docker to run special Node Exporters which will behave like servers that you can monitor. As a prerequisite, you will pull down the `bitnami/node-exporter` image from Docker Hub. You will use this image to make three instances of node exporters.

Your Task

1. Let's start by opening up a new terminal. From the top menu, go to **Terminal** and choose **New Terminal** to open a new terminal window.



2. Use the following `docker pull` command to pull down the `bitnami/node-exporter` image from Docker Hub that you will use to simulate 3 servers being monitored.

```
docker pull bitnami/node-exporter:latest
```

You are now ready to start the lab.

Step 1: Start the first node exporter

Now you need some server nodes to monitor. You will start up three node exporters listening on ports `9101`, `9102`, and `9103` respectively. In this step you will just start the first one and ensure that it is working correctly.

Your Task

1. Run the following command to create a docker network within which we will run all the docker instances.

```
docker network create monitor
```

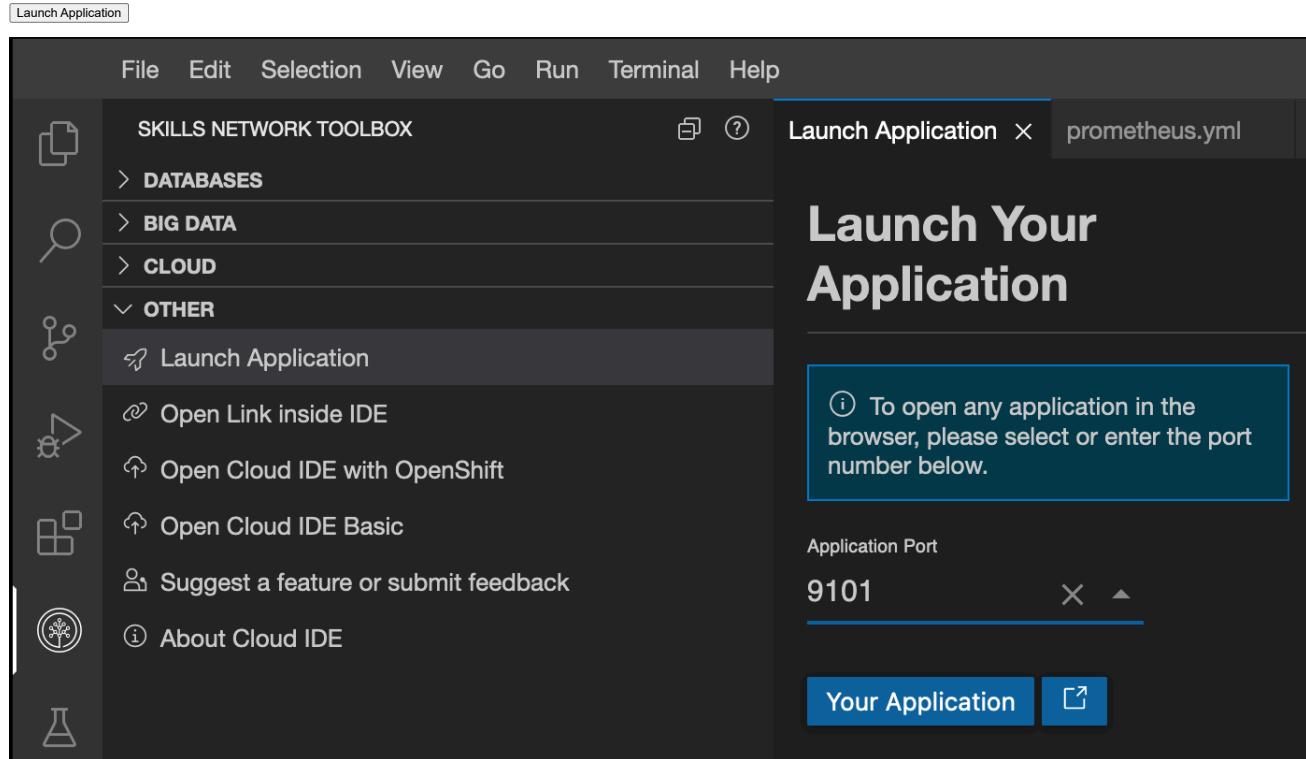
2. Run the following `docker run` command to start a node exporter instance listening at port `9101`

```
docker run -d -p 9101:9100 --name node-exporter1 --network monitor bitnami/node-exporter:latest
```

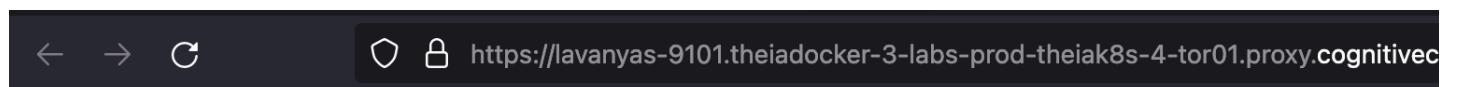
This will start a instance named node_exporter1 of node-exporter. The output should look like this:

```
theia@theiadocker-lavanyas:/home/project$ docker run -dp 9101:9100 --name node-exporter1 bitnami/node-exporter:5a601f5ff2c7c1e93f8a710082380183377d576ce05f290200f86eca140fb597
```

3. Next, check if the instance is running by either clicking on the **Skills Network Toolbox** and under **Others** choose **Launch Application**, enter the port number **9101** and click on the launch URL button. Or press the [**Launch Application**] button below:



4. The node exporter page opens up as below with a hyper link to **Metrics**. These are the metrics the Prometheus instance is going to monitor.



[Metrics](#)

5. Click on the **Metrics** link to have a glimpse of the metrics.

```
# HELP go_gc_duration_seconds A summary of the pause duration of garbage collection cycles.
# TYPE go_gc_duration_seconds summary
go_gc_duration_seconds{quantile="0"} 4.2085e-05
go_gc_duration_seconds{quantile="0.25"} 0.000158756
go_gc_duration_seconds{quantile="0.5"} 0.000198699
go_gc_duration_seconds{quantile="0.75"} 0.000231471
go_gc_duration_seconds{quantile="1"} 0.003538012
go_gc_duration_seconds_sum 0.344622043
go_gc_duration_seconds_count 1630
# HELP go_goroutines Number of goroutines that currently exist.
# TYPE go_goroutines gauge
go_goroutines 8
# HELP go_info Information about the Go environment.
# TYPE go_info gauge
go_info{version="go1.17.3"} 1
# HELP go_memstats_alloc_bytes Number of bytes allocated and still in use.
# TYPE go_memstats_alloc_bytes gauge
go_memstats_alloc_bytes 3.895568e+06
# HELP go_memstats_alloc_bytes_total Total number of bytes allocated, even if freed.
# TYPE go_memstats_alloc_bytes_total counter
go_memstats_alloc_bytes_total 2.846540048e+09
# HELP go_memstats_buck_hash_sys_bytes Number of bytes used by the profiling bucket hash table.
# TYPE go_memstats_buck_hash_sys_bytes gauge
go_memstats_buck_hash_sys_bytes 1.619023e+06
# HELP go_memstats_frees_total Total number of frees.
# TYPE go_memstats_frees_total counter
go_memstats_frees_total 2.0888216e+07
# HELP go_memstats_gc_cpu_fraction The fraction of this program's available CPU time used by the
# TYPE go_memstats_gc_cpu_fraction gauge
go_memstats_gc_cpu_fraction 8.542752870740738e-05
# HELP go_memstats_gc_sys_bytes Number of bytes used for garbage collection system metadata.
# TYPE go_memstats_gc_sys_bytes gauge
go_memstats_gc_sys_bytes 5.073568e+06
# HELP go_memstats_heap_alloc_bytes Number of heap bytes allocated and still in use.
# TYPE go_memstats_heap_alloc_bytes gauge
go_memstats_heap_alloc_bytes 3.895568e+06
# HELP go_memstats_heap_idle_bytes Number of heap bytes waiting to be used.
# TYPE go_memstats_heap_idle_bytes gauge
go_memstats_heap_idle_bytes 5.644288e+06
```

6. Copy the URL marked in red in the above image and save it in a notepad. You will use this URL to configure Prometheus to monitor the node exporter.

Step 2: Start two more node exporters

Now that you have one node exporter working, let's start two more so that Prometheus has three nodes to monitor in total. You will do this exactly the same way as you did the first node exporter except that you will change the port numbers to 9102 and 9103 respectively.

Your Task

1. In the terminal, now run the following commands to start two more instances of node exporter.

```
docker run -d -p 9102:9100 --name node-exporter2 --network monitor bitnami/node-exporter:latest
```

and

```
docker run -d -p 9103:9100 --name node-exporter3 --network monitor bitnami/node-exporter:latest
```

2. Now, check if all the instances of node exporter are running by using the docker ps command and grepping for node-exporter.

```
docker ps | grep node-exporter
```

Results

If everything started correctly, you should see output similar to the following coming back from the docker commands:

```
theia@theiadocker-lavanyas:/home/project$ docker run -dp 9101:9100 --name node-exporter1 bitnami/node-exporter:latest
5a601f5ff2c7c1e93f8a710082380183377d576ce05f290200f86eca140fb597
theia@theiadocker-lavanyas:/home/project$ docker run -dp 9102:9100 --name node-exporter2 bitnami/node-exporter:latest
06c0131da8db8756a3a003442f97a926c37378daf7d97790ef1fc4f6a4a0eeaca
theia@theiadocker-lavanyas:/home/project$ docker run -dp 9103:9100 --name node-exporter3 bitnami/node-exporter:latest
46e20a2bf0bd95d704587556921203f6f0aa134f0b3e9b3bbc230fdbdb8ffb51
theia@theiadocker-lavanyas:/home/project$ docker ps | grep node-exporter
46e20a2bf0bd    bitnami/node-exporter:latest    "/opt/bitnami/node-e..."    19 seconds ago    Up 18 seconds
:::9103->9100/tcp  node-exporter3
06c0131da8db    bitnami/node-exporter:latest    "/opt/bitnami/node-e..."    32 seconds ago    Up 31 seconds
:::9102->9100/tcp  node-exporter2
5a601f5ff2c7    bitnami/node-exporter:latest    "/opt/bitnami/node-e..."    19 minutes ago   Up 19 minutes
:::9101->9100/tcp  node-exporter1
theia@theiadocker-lavanyas:/home/project$
```

You are now ready to install and start Prometheus.

Step 3: Download, configure and run Prometheus

Now that we have some nodes to monitor, it's time to run Prometheus to monitor the nodes. In this step you will also create a custom configuration file to tell Prometheus what nodes to monitor. Then you will start Prometheus passing it the configuration file to use.

Your Task

1. First, pull the Prometheus into your local environment, by running the following `docker pull` command on your terminal.

```
docker pull bitnami/prometheus:latest
```

This may take a few seconds depending upon your internet connectivity

1. Create a file named `prometheus.yml` in the current directory. This is the file where you will configure the prometheus to monitor the node exporter instances.

```
touch prometheus.yml
```

2. From the explorer, navigate to `prometheus.yml` to edit the file or press the [Open `prometheus.yml` in IDE] button below:

[Open `prometheus.yml` in IDE](#)

3. Add the content below to the file. Update the URLs appropriately. The node exporter 1 URL will be the one that you copied into a notepad in the previous task. node exporter 2 will be the same URL with port number changed to 9102. node exporter 3 will be the same URL with port number changed to 9103.

```
# my global config
global:
  scrape_interval: 15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
  scrape_configs:
    - job_name: 'node'
      static_configs:
        - targets: ['node-exporter1:9100']
          labels:
            group: 'monitoring_node_ex1'
        - targets: ['node-exporter2:9100']
          labels:
            group: 'monitoring_node_ex2'
        - targets: ['node-exporter3:9100']
          labels:
            group: 'monitoring_node_ex3'
```

A sample configuration can be seen in the image below.

```

prometheus.yml
1  # my global config
2  global:
3    scrape_interval: 15s # Set the scrape interval to every 15 seconds. Default is every 1
4
5  scrape_configs:
6    - job_name: 'node'
7
8      static_configs:
9        - targets: ['lavanyas-9101.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitive
10          labels:
11            group: 'monitoring_node_ex1'
12        - targets: ['lavanyas-9102.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitive
13          labels:
14            group: 'monitoring_node_ex2'
15        - targets: ['lavanyas-9103.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitive
16          labels:
17            group: 'monitoring_node_ex3'

```

4. Now you can launch the prometheus monitor by passing this yaml file as a parameter.

```

docker run --rm --name prometheus -p 9090:9090 --network monitor \
-v $(pwd)/prometheus.yml:/opt/bitnami/prometheus/conf/prometheus.yml \
bitnami/prometheus:latest

```

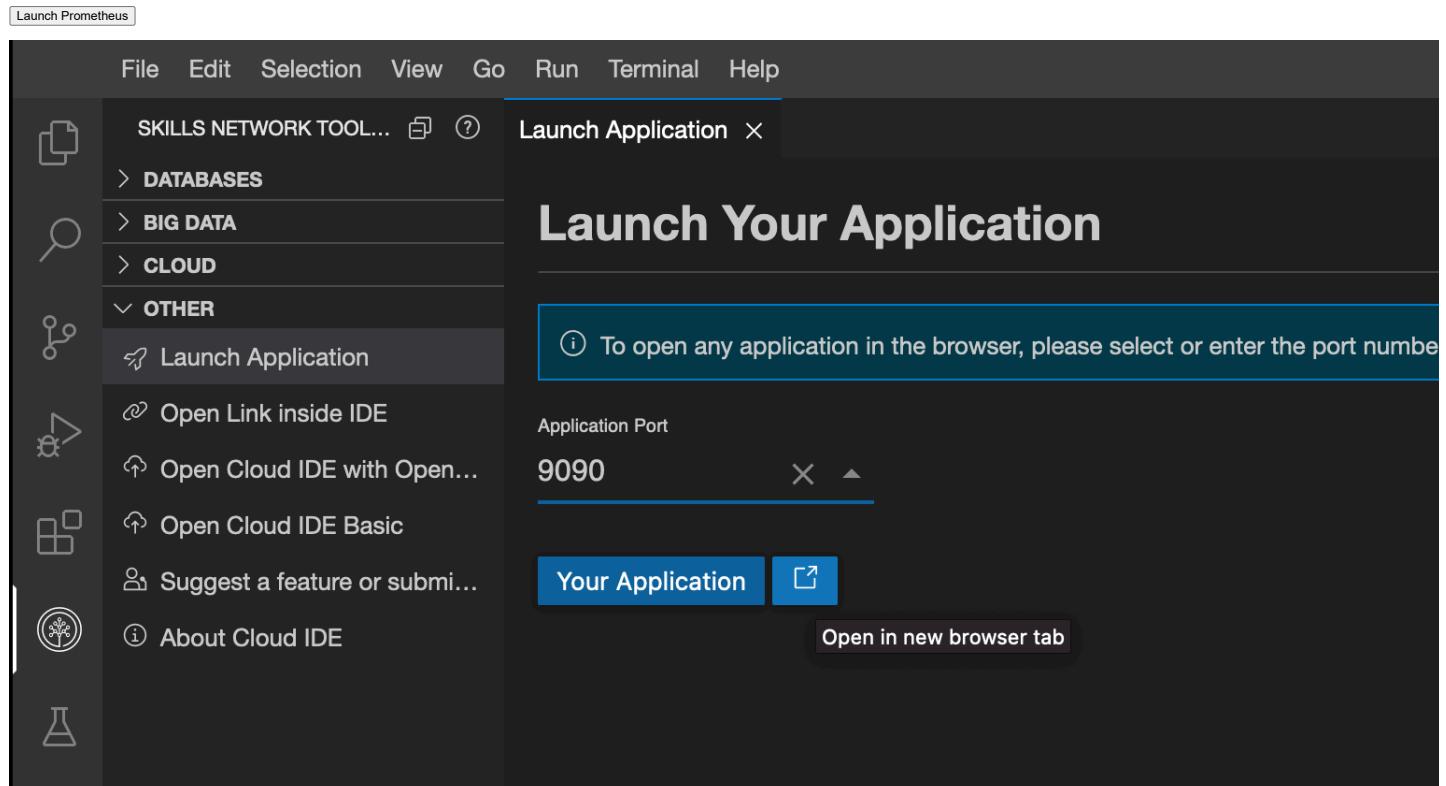
Results

You should see the Prometheus logs scroll up the screen which indicate that the monitoring has started.

Step 4: Open the Prometheus UI

In this step you will launch the Prometheus web UI and navigate to the page where you can start executing queries.

1. Open the Prometheus web UI by click on the **Skills Network Toolbox** and under **Others** choose **Launch Application**, enter the port number **9090** and click on the launch URL button or press the Launch Prometheus button below to launch in an external browser:



2. The Prometheus application UI opens up by default in the graph end-point.

The screenshot shows the Prometheus application interface. At the top, there is a navigation bar with icons for back, forward, and refresh, followed by a URL field displaying `https://lavanyas-9090.theiadocker-3-labs-prod-theiak8s-4-tor01.proxy.c`. Below the URL is the Prometheus logo and the word "Prometheus". To the right of the logo are links for "Alerts", "Graph", "Status", and "Help". A dropdown menu icon is also present. Below the navigation bar, there is a row of checkboxes with the following options: "Use local time" (unchecked), "Enable query history" (checked), "Enable autocomplete" (checked), "Enable highlighting" (checked), and another checkbox which is partially visible and appears to be checked. To the left of these checkboxes is a magnifying glass icon. Next to the icon is a text input field with the placeholder text "Expression (press Shift+Enter for newlines)". Below the input field are two tabs: "Table" (grayed out) and "Graph" (highlighted in blue). Underneath the tabs is a search bar with a left arrow, a central input field containing "Evaluation time", and a right arrow. Below the search bar, a message says "No data queried yet". At the bottom left of the main content area is a blue button labeled "Add Panel".

3. Next, click on **Status** on the menu and choose **Targets** to see which targets are being monitored.

The screenshot shows the Prometheus web interface. At the top, there's a navigation bar with icons for back, forward, and refresh, followed by a URL field showing `https://lavanyas-9090.theiadocker-3-labs-prod-theiak8s-4-tor01.proxy.cognit...`. Below the URL is the Prometheus logo and the word "Prometheus". To the right of the logo are links for "Alerts", "Graph", "Status ▾", and "Help". A dropdown menu is open over the "Status" link, listing options: "Runtime & Build Information", "TSDB Status", "Command-Line Flags", "Configuration", "Rules", "Targets", and "Service Discovery". On the left side of the main content area, there are checkboxes for "Use local time" and "Enable query his...", a search bar with placeholder text "Expression (press Shift+Enter for r...", and tabs for "Table" and "Graph". Below these is a date range selector with "Evaluation time" and arrows, and a message "No data queried yet". At the bottom left is a blue "Add Panel" button.

4. View the status of all the three node exporters.

The screenshot shows the "Targets" page in the Prometheus interface. At the top, there's a navigation bar with the Prometheus logo and the word "Prometheus", followed by links for "Alerts", "Graph", "Status ▾", and "Help". Below the navigation is a large heading "Targets". Underneath the heading are buttons for "All", "Unhealthy", and "Collapse All", with "All" currently selected. To the right is a search bar and a "Filter by endpoint or labels" input field. The main content is a table titled "node (3/3 up)". The table has columns for "Endpoint", "State", and "Labels". There are three rows, each representing a node exporter:

| Endpoint | State | Labels |
|--|-------|---|
| <code>http://lavanyas-9101.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/metrics</code> | UP | <code>group="monitoring_node_ex1"</code> <code>instance="lavanyas-9101.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/metrics"</code> <code>job="node"</code> |
| <code>http://lavanyas-9102.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/metrics</code> | UP | <code>group="monitoring_node_ex2"</code> <code>instance="lavanyas-9102.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/metrics"</code> <code>job="node"</code> |
| <code>http://lavanyas-9103.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/metrics</code> | UP | <code>group="monitoring_node_ex3"</code> <code>instance="lavanyas-9103.theiadocker-2-labs-prod-theiak8s-4-tor01.proxy.cognitiveclass.ai/metrics"</code> <code>job="node"</code> |

The screenshot shows the Prometheus Graph interface. At the top, there are navigation icons (back, forward, search) and a URL bar showing <https://lavanyas-9090.theiadocker-3-labs-prod-theiak8s-4-tor01.proxy.co>. Below the URL is the Prometheus logo and the word "Prometheus". The "Graph" tab is highlighted with a red box. Other tabs include "Alerts", "Status", and "Help".

Targets

Filter by endpoint or labels

node (3/3 up) [show less](#)

| Endpoint | State | Labels | Last check |
|---|-------|---|------------|
| http://localhost:9101/metrics | UP | group="monitoring_node_ex1" instance="localhost:9101" job="node" | 4.4 |
| http://localhost:9102/metrics | UP | group="monitoring_node_ex2" instance="localhost:9102" job="node" | 7.13 |
| http://localhost:9103/metrics | UP | group="monitoring_node_ex3" instance="localhost:9103" job="node" | 3.9 |

You are now ready to execute queries.

Step 5: Execute your first query

You are now ready to execute your first query. Query for the total CPU second of the nodes by pasting the following query and executing it. It will show the graph as given in the image. You can observe the details for each instance by hovering the mouse over that instance.

Your Task

- Pasting the following query and executing it. It will show the graph as given in the image. You can observe the details for each instance by hovering the mouse over that instance.

```
node_cpu_seconds_total
```

 Use local time Enable query history Enable autocomplete Enable highlighting

node_cpu_seconds_total

[Table](#)[Graph](#)

1h



End time



Res. (s)



Show Examples

60.00k

50.00k

40.00k

30.00k

20.00k

2022-09-27 00:00:00

node_cpu_seconds_total{cpu:3,group:monitoring,instance:lava,job:node,mode:idle}

Series:

node_cpu_seconds_total{cpu:3,group:monitoring,instance:lava,job:node,mode:idle}

2. Next on [Table](#) and query the cpu second for all the targets in tabular format.



Use local time Enable query history Enable autocomplete Enable highlighting

node_cpu_seconds_total

Table

Graph

< Evaluation time >

3. Now, filter the query to get the details for only one instance using the following query.

```
node cpu seconds total{group="monitoring node ex?"}
```



Use local time Enable query history Enable autocomplete Enable highlighting

```
node_cpu_seconds_total{group="monitoring_node_ex2"}
```

Table

Graph

< Evaluation time >

4. Finally, query for the connections each node has using this query:

node ips connections total



Use local time Enable query history Enable autocomplete Enable highlighting



node_ipvs_connections_total|

Table

Graph



Evaluation time



```
node_ipvs_connections_total{group="monitoring_node_ex1", instance="lavanyas-9101.theiadocker-2-labs-prod-theia"}  
node_ipvs_connections_total{group="monitoring_node_ex2", instance="lavanyas-9102.theiadocker-2-labs-prod-theia"}  
node_ipvs_connections_total{group="monitoring_node_ex3", instance="lavanyas-9103.theiadocker-2-labs-prod-theia"}
```

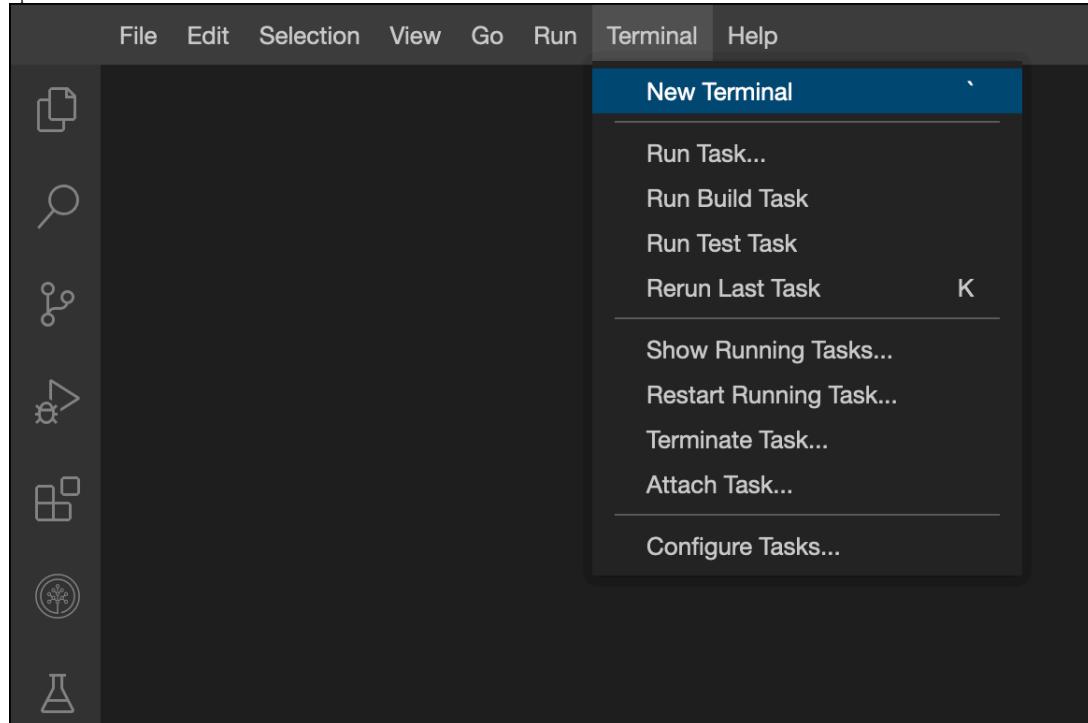
[Add Panel](#)

Step 6: Stop and observe

In this step we will stop one of the node exporter instances and see how that is reflected in the Prometheus console.

Your Task

1. Open a new terminal.



2. Stop the node-exporter1 instance by running the following `docker stop` command and then switch back to old terminal in which prometheus is running.

```
docker stop node-exporter1
```

3. Now go back to the prometheus UI on your browser and check the targets by selecting the menu item **Status > Targets**

Results

You should now see that one of the node exporters is being monitored is down. *Note: You may have to refresh your browser to see the new status*

The screenshot shows the Prometheus Targets page. At the top, there are navigation icons and a URL bar showing <https://lavanyas-9090.theiadocker-3-labs-prod-theiak8s-4-tor01.proxy.co>. Below the header, there's a navigation bar with links for Prometheus, Alerts, Graph, Status, and Help. The main title is "Targets". A search bar and a filter input field ("Filter by endpoint or labels") are also present.

The main content area displays a table of targets:

| Endpoint | State | Labels |
|---|-------|---|
| http://localhost:9101/metrics | DOWN | group="monitoring_node_ex1" instance="localhost:9101" job="node" |
| http://localhost:9102/metrics | UP | group="monitoring_node_ex2" instance="localhost:9102" job="node" |
| http://localhost:9103/metrics | UP | group="monitoring_node_ex3" instance="localhost:9103" job="node" |

Step 7: Start Grafana

1. Run the following command to pull Grafana from docker and run it. We will run it in the same network as node monitor and prometheus.

```
docker run --name=grafana -dp 3000:3000 --network monitor grafana/grafana
```

2. Now press Launch Application button below to open the Grafana browser by connection.

[Launch Grafana](#)

3. Log in with username as admin and password as admin. You may choose to set a new password. This will take you to the Grafana homepage.

If you set a new password, it will only last for that docker instance. When you restart you need to reset it.

The screenshot shows the Grafana Home page. At the top left is the 'General / Home' icon. The top right features a search bar, a star icon, and a gear icon. Below the header, the main content area includes:

- Welcome to Grafana**: A large title.
- Basic**: A section with a yellow gear icon and text: "The steps below will guide you to quickly finish setting up your Grafana installation." It also includes a "TUTORIAL" link to "DATA SOURCE AND DASHBOARDS" and "Grafana fundamentals".
- Data Sources**: A panel titled "Add your first data source" with a database icon. It includes a "Learn how in the docs" link.
- COMPLETE**: A panel titled "Create your first dashboard" with a dashboard icon. It includes a "Learn how in the docs" link.
- Dashboards**: A sidebar with sections for "Starred dashboards" (empty), "Recently viewed dashboards" (empty), and "Node Exporter Full".
- Latest from the blog**: Two blog posts:
 - Grafana 9.2 release: New Grafana panel help options, Grafana auth updates, simplified variable editor for Grafana Loki, and more!** (Oct 10)
 - How to land a job as a junior engineer at Grafana Labs** (Oct 06)

4. Click on Data Sources to add your first data source.

5. Choose Prometheus from the list of options available.

Note: In this lab, we are focusing on using Prometheus to monitor the node exporters and visualize the metrics and graphs. Although Step 7 mentions starting Grafana, the lab is designed to demonstrate the capabilities of Prometheus in monitoring the node exporters. Therefore, we are not proceeding with configuring Grafana for this particular lab.

Not Secure — 778dc6371c67d206ba9ae5c-0000.

Add data source

Choose a data source type

Filter by name or type

Time series databases

- Prometheus**
Open source time series database & alerting
Core
- Graphite**
Open source time series database
Core
- InfluxDB**
Open source time series database
Core
- OpenTSDB**
Open source time series database
Core

Logging & document databases

- Loki**
Like Prometheus but for logs. OSS logging solution from Grafana Labs
Core
- Elasticsearch**
Open source logging & analytics database
Core

Conclusion

Congratulations! You have completed this lab on Prometheus. You are now well on your way to monitoring your own applications to ensure they are running properly.

Next Steps

Your next challenge is to set up Prometheus in your development environment to monitor your applications. You can use some of the queries you learned in this lab to check on the health and performance of your applications.

Author(s)

Lavanya T S

Other Contributor(s)

